

HPCSAM CASE STUDY

TITLE: Build a two node disk less HPC cluster using openHPC with xCAT, openLDAP, slurm, HPL benchmarking, ganglia and document the results.

NAME	PRN NO.
SUMIT KUMAR	220940127021
VARSHARANI MAHAVIR DANOLE	220940127022
VISHAL BHAUSAHEB MAGAR	220940127023
WAGHMARE TEJAS ARUN	220940127024
AMOL MURLIDHAR DEOGADE	220940127031

XCAT Configuration

1. Launch VM with configuration:

➤ Master node :

- RAM - 6GB
- Processor -2 & Core - 2
- HDD-50GB
- Network Adaptors - 2 (1 NAT, 1 Custom Host(disable DHCP))

➤ Client node1 & node2 :

- Ram - 4GB
- Processor - 2 & Core -2
- HDD -20GB
- Network Adaptor- 1(Custom Host as Master's Custom Host)

2. Provide manual IP & Check IP

```
# nmtui
```

```
#ip a
```

3. Stop and Disable Firewall

```
#systemctl stop firewalld
```

```
#systemctl disable firewalld
```

```
#systemctl status firewalld
```

4. Disable Selinux

```
#vim /etc/selinux/config
```

```
SELINUX= disabled
```

```
#sestatus
```

5. Install Utilities

```
#yum install yum-utils
```

6. Download xCAT core/dependencies repository and Install xCAT

```
#wget -P /etc/yum.repos.d https://xcat.org/files/xcat/repos/yum/latest/xcat-core/xcat-core.repo --no-check-certificate
```

```
#wget -P /etc/yum.repos.d https://xcat.org/files/xcat/repos/yum/xcat-dep/rh7/x86\_64/xcat-dep.repo --no-check-certificate
```

```
#yum install xCAT
```

7. Set System Environment

```
#. /etc/profile.d/xcat.sh
```

```
#echo $PATH
```

8. Set Interface and IP

```
#chdef -t site dhcpinterfaces="ens34"
```

```
#chdef -t site master="192.168.100.11"
```

```
#tabdump site | grep master
```

```
#tabdump site | grep dhcp
```

9. Copy ISO file

```
#dd if=/dev/sr0 of=/root/hpcsas2os7.iso  
#copycds /root/hpcsas2os7.iso  
#lsdef -t osimage
```

10. Generate Image

```
#genimage hpcsas2os7.9-x86_64-netboot-compute
```

11. Make Directory and add files

```
#mkdir -p /install/custom/netboot  
#chdef -t osimage hpcsas2os7.9-x86_64-netboot-compute synclist="/install/custom/netboot/compute.synclist"  
  
#echo "/etc/passwd -> /etc/passwd" > /install/custom/netboot/compute.synclist  
#echo "/etc/shadow -> /etc/shadow" >> /install/custom/netboot/compute.synclist  
#echo "/etc/gshadow -> /etc/gshadow" >> /install/custom/netboot/compute.synclist  
#echo "/etc/group -> /etc/group" >> /install/custom/netboot/compute.synclist  
#echo "/etc/hosts -> /etc/hosts" >> /install/custom/netboot/compute.synclist
```

12. Pack Image

```
#packimage hpcsas2os7.9-x86_64-netboot-compute
```

13. Copy MAC address of node1 & node2 and Assign

```
#mkdef -t node node1 groups=compute,all ip=192.168.100.17 mac=00:0C:29:AF:60:A7 netboot=xnba  
#mkdef -t node node2 groups=compute,all ip=192.168.100.20 mac=00:0C:29:26:3D:D7 netboot=xnba  
#lsdef node  
#chdef -t group compute provmethod=hpcsas2os7.9-x86_64-netboot-compute  
#chdef -t site domain=xcat.in
```

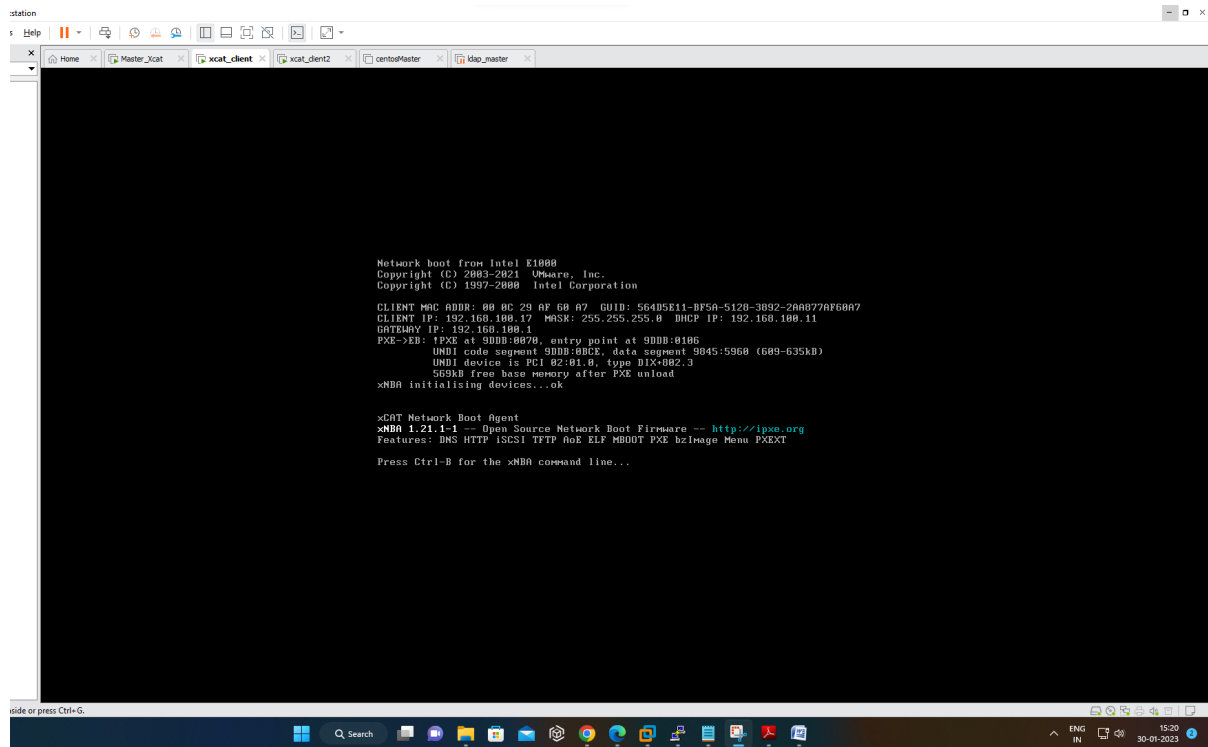
14. Create

```
#makehosts  
#makenetworks  
#makedhcp -n  
#makedns -n  
#nodeset compute osimage=hpcsas2os7.9-x86_64-netboot-compute
```

15. Restart DHCPD

```
#systemctl restart dhcpd
```

16. Restart node1 & node2-

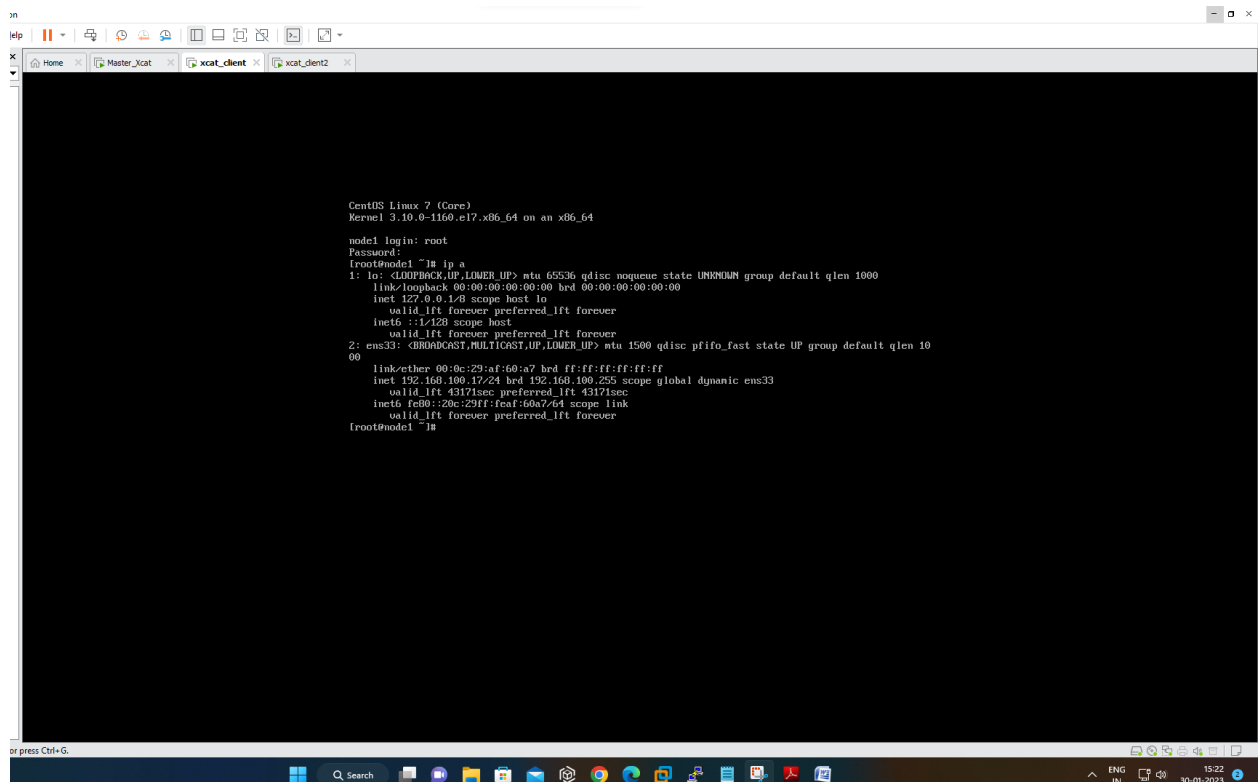


The screenshot shows a terminal window titled 'station' with a Windows taskbar at the bottom. The terminal displays the following text:

```
Network boot from Intel E1000
Copyright (C) 2003-2021 VMware, Inc.
Copyright (C) 1997-2000 Intel Corporation

CLIENT MAC ADDR: 00 0C 29 AF 60 A7 GUID: 564D5E11-BF5A-5120-3092-2A0077AF60A7
CLIENT IP: 192.168.100.17 MASK: 255.255.255.0 DHCP IP: 192.168.100.11
GATEWAY IP: 192.168.100.1
PXE->EB: PXE at 9000:0070, entry point at 9000:0100
        UNDI code segment 9000:0000, data segment 9045:5960 (609-635KB)
        UNDI device is PCI 02:01:0, type DIX-002.3
        569KB free base memory after PXE unload
xNBA initialising devices...ok

xCAT Network Boot Agent
xNBA 1.21.1-1 -- Open Source Network Boot Firmware -- http://ipxe.org
Features: DNS HTTP iSCSI TFTP AoE ELF MBOOT PXE bzImage Menu PXEXT
Press Ctrl-B for the xNBA command line...
```



The screenshot shows a terminal window titled 'jn' with a Windows taskbar at the bottom. The terminal displays the following text:

```
CentOS Linux 7 (Core)
Kernel 3.10.0-1160.el7.x86_64 on an x86_64

node1 login: root
Password:
[root@node1 ~]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 00:0c:29:af:60:a7 brd ff:ff:ff:ff:ff:ff
    inet 192.168.100.17/24 brd 192.168.100.255 scope global dynamic ens3
        valid_lft 43171sec preferred_lft 43171sec
    inet6 fe80::20c:29ff:feaf:60a7/64 scope link
        valid_lft forever preferred_lft forever
[root@node1 ~]#
```

```
CentOS Linux 7 (Core)
Kernel 3.10.0-1160.el7.x86_64 on an x86_64

node2 login: [ 35.419737] 0021q: 002.1Q VLAN Support v1.0
[ 35.419763] 0021q: adding VLAN 0 to HW filter on device ens33

CentOS Linux 7 (Core)
Kernel 3.10.0-1160.el7.x86_64 on an x86_64

node2 login: root
Password:
[root@node2 ~]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 00:0c:29:26:3d:47 brd ff:ff:ff:ff:ff:ff
    inet 192.168.100.20/24 brd 192.168.100.255 scope global dynamic ens33
        valid_lft 43169sec preferred_lft 43169sec
    inet6 fe80::20c:29ff:fe26:3d47/64 scope link
        valid_lft forever preferred_lft forever
[root@node2 ~]#
```

LDAP Configuration

➤ On Master Node:-

- **Install OpenLDAP**

```
#yum -y install openldap-servers openldap-clients
#cp /usr/share/openldap-servers/DB_CONFIG.example /var/lib/ldap/DB_CONFIG
#chown ldap. /var/lib/ldap/DB_CONFIG
#systemctl start slapd
#systemctl enable slapd
# generate encrypted password
#save this password to a file
# password from user
#vi chrootpw.ldif
```


specify the password generated above for "olcRootPW" section

dn: olcDatabase={0}config,cn=config

changetype: modify

add: olcRootPW

olcRootPW: {SSHA}xxxxxxxxxxxxxxxxxxxxxxxxxxxx



```
#ldapadd -Y EXTERNAL -H ldapi:/// -f chrootpw.ldif
#ldapadd -Y EXTERNAL -H ldapi:/// -f /etc/openldap/schema/cosine.ldif
#ldapadd -Y EXTERNAL -H ldapi:/// -f /etc/openldap/schema/nis.ldif
#ldapadd -Y EXTERNAL -H ldapi:/// -f /etc/openldap/schema/inetorgperson.ldif
```

- **generate directory manager's password**

```
#slappasswd
```

```
#vi chdomain.ldif
```

```
# replace to your own domain name for "dc=***,dc=***" section
# specify the password generated above for "olcRootPW" section
dn: olcDatabase={1}monitor,cn=config
changetype: modify
replace: olcAccess
olcAccess: {0}to * by dn.base="gidNumber=0+uidNumber=0,cn=peercred,cn=external,cn=auth"
read by dn.base="cn=Manager,dc=casestudy,dc=in" read by * none
dn: olcDatabase={2}hdb,cn=config
changetype: modify
replace: olcSuffix
olcSuffix: dc=casestudy,dc=in
dn: olcDatabase={2}hdb,cn=config
changetype: modify
replace: olcRootDN
olcRootDN: cn=Manager,dc=casestudy,dc=in
dn: olcDatabase={2}hdb,cn=config
changetype: modify
add: olcRootPW
olcRootPW: {SSHA}xxxxxxxxxxxxxxxxxxxxxxxxxxxx
dn: olcDatabase={2}hdb,cn=config
changetype: modify
add: olcAccess
olcAccess: {0}to attrs=userPassword,shadowLastChange by
dn="cn=Manager,dc=casestudy,dc=in" write by anonymous auth by self write by * none
olcAccess: {1}to dn.base="" by * read
olcAccess: {2}to * by dn="cn=Manager,dc=casestudy,dc=in" write by * read
```

```
#ldapmodify -Y EXTERNAL -H ldapi:/// -f chdomain.ldif
```



```
File Edit View Search Terminal Help
# create new
# replace to your own domain name for "dc=*,dc=*" section
dn: uid=hpcs2,ou=People,dc=casestudy,dc=in
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: hpcs2
sn: hpcs2
userPassword: {SSHA}CpBm1dWt7RAG/Wj5r5d3gtvu3070g
loginShell: /bin/bash
uidNumber: 1003
gidNumber: 1003
homeDirectory: /home/hpcs2

dn: cn=hpcs2,ou=group,dc=casestudy,dc=in
objectClass: posixGroup
cn: hpcs2
gidNumber: 1003
memberUid: hpcs2

~ldapuser.ldif* 19L, 487C

root@master1:~#
```

#ldapadd -x -D cn=Manager,dc=casestudy,dc=in -W -f ldapuser.ldif

#systemctl start slapd

#systemctl status slapd

```
[root@master1 ~]# systemctl status slapd
● slapd.service - OpenLDAP Server Daemon
   Loaded: loaded (/usr/lib/systemd/system/slapd.service; enabled; vendor preset: disabled)
   Active: active (running) since Sun 2023-01-29 19:08:20 IST; 21h ago
     Docs: man:slapd
           man:slapd-config
           man:slapd-hdb
           man:slapd-mdb
           file:///usr/share/doc/openldap-servers/guide.html
   Process: 21238 ExecStart=/usr/sbin/slapd -u ldap -h ${SLAPD_URLS} $SLAPD_OPTIONS (code=exited, status=0/SUCCESS)
   Process: 21200 ExecStartPre=/usr/libexec/openldap/check-config.sh (code=exited, status=0/SUCCESS)
   Main PID: 21241 (slapd)
      Tasks: 4
     Memory: 32.0K
    CGroup: /system.slice/slapd.service
            └─21241 /usr/sbin/slapd -u ldap -h ldapi:/// ldap:///

Jan 29 19:08:32 master1 slapd[21241]: conn=1000 op=1 SEARCH RESULT tag=101 err=0 nentries=10 text=
Jan 29 19:08:32 master1 slapd[21241]: conn=1000 op=2 UNBIND
Jan 29 19:08:32 master1 slapd[21241]: conn=1000 fd=11 closed
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 fd=11 ACCEPT from IP=192.168.100.11:60158 (IP=0.0.0.0:389)
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 op=0 BIND dn="" method=128
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 op=0 RESULT tag=97 err=0 text=
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 op=1 SRCH base="dc=casestudy,dc=in" scope=2 deref=0 filter="(objectClass=*)"
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 op=1 SEARCH RESULT tag=101 err=0 nentries=10 text=
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 op=2 UNBIND
Jan 30 16:29:08 master1 slapd[21241]: conn=1001 fd=11 closed
[root@master1 ~]#
```

• On clientNodes

Install OpenLDAP Client.

#yum --installroot=/install/netboot/hpcs2os7.9/x86_64/compute/rootimg install openldap-clients

#exports CHROOT=/install/netboot/hpcs2os7.9/x86_64/compute/rootimg

#chroot \$CHROOT

#authconfig --enableldap --enableldapauth --ldapsrvr=master --ldapbasedn="dc=casestudy,dc=in" --

enablemkhomedir --update

#systemctl restart nslcd

#systemctl enable nslcd

#systemctl status nslcd

```
[root@node1 ~]# systemctl status nslcd
● nslcd.service - Naming services LDAP client daemon.
   Loaded: loaded (/usr/lib/systemd/system/nslcd.service; enabled; vendor preset: disabled)
   Active: active (running) since Mon 2023-01-30 20:13:59 IST; 1min 42s ago
     Docs: man:nslcd(8)
           man:nslcd.conf(5)
  Process: 1371 ExecStart=/usr/sbin/nslcd (code=exited, status=0/SUCCESS)
 Main PID: 1484 (nslcd)
    Tasks: 6
   CGroup: /system.slice/nslcd.service
           └─1484 /usr/sbin/nslcd

Jan 30 20:14:30 localhost.localdomain nslcd[1484]: [0bd78f] <group/member="ro...
Jan 30 20:14:30 localhost.localdomain nslcd[1484]: [0bd78f] <group/member="ro...
Jan 30 20:14:31 localhost.localdomain nslcd[1484]: [3ea438] <group/member="ge...
Jan 30 20:14:31 localhost.localdomain nslcd[1484]: [3ea438] <group/member="ge...
Jan 30 20:15:04 localhost.localdomain nslcd[1484]: [55585c] <group/member="ro...
Jan 30 20:15:04 localhost.localdomain nslcd[1484]: [55585c] <group/member="ro...
Jan 30 20:15:04 localhost.localdomain nslcd[1484]: [55585c] <group/member="ro...
Jan 30 20:15:17 nodel nslcd[1484]: [a64e2a] <group/member="root"> failed to...ed
Jan 30 20:15:17 nodel nslcd[1484]: [a64e2a] <group/member="root"> no availa...ed
Jan 30 20:15:17 nodel nslcd[1484]: [a64e2a] <group/member="root"> no availa...ed
Hint: Some lines were ellipsized, use -l to show in full.
[root@node1 ~]#
```

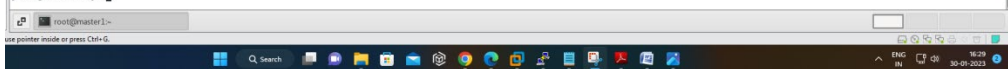
#ldapsearch -x

```
# hpcsa2, People, casestudy, in
dn: uid=hpcsa2,ou=People,dc=casestudy,dc=in
objectClass: inetOrgPerson
objectClass: posixAccount
objectClass: shadowAccount
cn: hpcsa2
sn: Linux
userPassword:: e1NTSEF9038CV25sZFVtVHpkUkcV2pTclNKM2dsdKZ35jA3UWc=
loginShell: /bin/bash
uidNumber: 1003
gidNumber: 1003
homeDirectory: /home/hpcsa2
uid: hpcsa2

# hpcsa2, Group, casestudy, in
dn: cn=hpcsa2,ou=group,dc=casestudy,dc=in
objectClass: posixGroup
cn: hpcsa2
gidNumber: 1003
memberUid: hpcsa2

# search result
search: 2
result: 0 Success

# numResponses: 11
# numEntries: 10
[root@master1 ~]#
```



```
rpc:x:32:32:Rpcbind Daemon:/var/lib/rpcbind:/sbin/nologin
sane:x:996:993:SANE scanner daemon user:/usr/share/sane:/sbin/nologin
sasauth:x:995:76:Sasauthd user:/run/sasauthd:/sbin/nologin
abrt:x:173:173::/etc/abrt:/sbin/nologin
setroubleshoot:x:994:991::/var/lib/setroubleshoot:/sbin/nologin
rtkit:x:172:172:RealtimeKit:/proc:/sbin/nologin
pulse:x:171:171:PulseAudio System Daemon:/var/run/pulse:/sbin/nologin
chrony:x:993:993::/var/lib/chrony:/sbin/nologin
unbound:x:992:987:Unbound DNS resolver:/etc/unbound:/sbin/nologin
radvd:x:75:75:radvd user:/:/sbin/nologin
tss:x:59:59:Account used by the trousers package to sandbox the tcsd daemon:/dev/null:/sbin/nologin
usbmuxd:x:113:113:usbmuxd user:/:/sbin/nologin
geoclue:x:991:985:User for geoclue:/var/lib/geoclue:/sbin/nologin
qemu:x:107:107:qemu user:/:/sbin/nologin
gluster:x:990:984:GlusterFS daemons:/run/gluster:/sbin/nologin
gdm:x:42:42::/var/lib/gdm:/sbin/nologin
rpcuser:x:29:29:RPC Service User:/var/lib/nfs:/sbin/nologin
nfsnobody:x:65534:65534:Anonymous NFS User:/var/lib/nfs:/sbin/nologin
gnome-initial-setup:x:989:983::/run/gnome-initial-setup:/sbin/nologin
sshd:x:74:74:Privilege-separated SSH:/var/empty/ssh:/sbin/nologin
avahi:x:70:70:Avahi mDNS/DNS-SD Stack:/var/run/avahi-daemon:/sbin/nologin
postfix:x:89:89::/var/spool/postfix:/sbin/nologin
ntp:x:38:38::/etc/ntp:/sbin/nologin
tcpdump:x:72:72::/sbin/nologin
hpcsa2x:1000:1000:hpcsa2:/home/hpcsa2/bin/bash
ldap:x:55:55:OpenLDAP server:/var/lib/ldap:/sbin/nologin
test:x:1001:1001::/home/test:/bin/bash
nscd:x:28:28:NSCD Daemon:/:/sbin/nologin
nslcd:x:65:55:LDAP Client User:/:/sbin/nologin
test:x:1001:1001:test:/home/test:/bin/bash
```

Slurm Configuration

- **On Master Node**

```
#wget https://download.schedmd.com/slurm/slurm-22.05.7.tar.bz2
#rpmbuild slurm-22.0507.tar.bz2
#yum install -y mariadb-server mariadb-devel epel-release munge munge-libs munge-devel rpm-build python3
perl-ExtUtils-Install gcc
#yum install openssl openssl-devel pam-devel numactl numactl-devel hwloc hwloc-devel lua lua-devel
readline-devel rrdtool-devel ncurses-devel man2html libibmad libibumad -y
#/usr/sbin/create-munge-key -r
#chown munge:munge /etc/munge
#chmod 400 /etc/munge/munge.key
#cd /root/rpmbuild/RPMS/x86_64/
#yum install slurm*
#export SLURMUSER=900
#groupadd -g $SLURMUSER slurm
#useradd -m -c "SLURM workload manager" -d /var/lib/slurm -u $SLURMUSER -g slurm -s /bin/bash slurm
#cp /etc/slurm/slurm.conf.example /etc/slurm/slurm.conf
#vi /etc/slurm/slurm.conf

    clusturname=oxygen
```

```
##
# slurm.conf file generated by configurator.html.
# Put this file on all nodes of your cluster.
# See the slurm.conf man page for more information.
#
ClusterName=oxygen
SlurmctldHost=master1
#SlurmctldHost=
#
#DisableRootJobs=NO
#EnforcePartLimits=NO
```

```
#mkdir -p /var/share/slurm/ctld
#vi /etc/slurm/slurm.conf

    stateSaveLocation=/var/share/slurm/ctld
    SlurmSpoolDir=/var/share/slurm/d
```

```
SlurmctldPidFile=/var/run/slurmctld.pid
SlurmctldPort=6817
SlurmdPidFile=/var/run/slurmd.pid
SlurmdPort=6818
SlurmdSpoolDir=/var/share/slurm/d
SlurmUser=slurm
#SlurmUser=root
#SrunEpilog=
#SrunProlog=
StateSaveLocation=/var/share/slurm/ctld
SwitchType=switch/none
#TaskEpilog=
TaskPlugin=task/affinity
#TaskProlog=
#TopologyPlugin=topology/tree
#TmpFS=/tmp
#TrackWcKey=no
#TreeWidth=
#UnkillableStepProgram=
#UsePAM=0
#
#
```

```
#chown -R slurm:slurm /var/share/slurm
```

```
#touch /var/log/slurmd.log
```

```
#systemctl start slurmd
```

```
#systemctl status slurmd
```

```
#systemctl enable slurmd
```

```
[root@master1 ~]# systemctl status slurmd.service
● slurmd.service - Slurm controller daemon
   Loaded: loaded (/usr/lib/systemd/system/slurmd.service; enabled; vendor preset: disabled)
   Active: active (running) since Mon 2023-01-30 11:56:44 IST; 4h 47min ago
     Main PID: 42651 (slurmd)
        Tasks: 10
       Memory: 3.2M
      CGroup: /system.slice/slurmd.service
              └─42651 /usr/sbin/slurmd -D -s
                 42658 slurmd: slurmdscriptd

Jan 30 12:08:10 master1 slurmd[42651]: slurmd: sched: slurm_rpc allocate_resources JobId=7 NodeList=(null) usec=1961
Jan 30 12:09:08 master1 slurmd[42651]: slurmd: slurm_rpc_kill_job: REQUEST_KILL_JOB JobId=4 uid 0
Jan 30 12:09:48 master1 slurmd[42651]: slurmd: sched: slurm_rpc allocate_resources JobId=8 NodeList=(null) usec=5383
Jan 30 15:49:26 master1 slurmd[42651]: slurmd: Node node2 now responding
Jan 30 15:50:25 master1 slurmd[42651]: slurmd: cleanup completing: JobId=4 completion process took 13487 seconds
Jan 30 15:50:25 master1 slurmd[42651]: slurmd: sched/backfill: start_job: Started JobId=7 in debug on node[1-2]
Jan 30 15:50:35 master1 slurmd[42651]: slurmd: job complete: JobId=7 done
Jan 30 15:50:36 master1 slurmd[42651]: slurmd: sched: Allocate JobId=8 NodeList=node[1-2] #CPUs=2 Partition=debug
Jan 30 15:50:46 master1 slurmd[42651]: slurmd: job complete: JobId=8 done
Jan 30 15:51:45 master1 slurmd[42651]: slurmd: error: Nodes node2 not responding
[root@master1 ~]#
```

```
#systemctl status munge
```

```
#systemctl start munge
```

```
#systemctl enable munge
```

```
[root@master1 ~]# systemctl status munge.service
● munge.service - MUNGE authentication service
   Loaded: loaded (/usr/lib/systemd/system/munge.service; enabled; vendor preset: disabled)
   Active: active (running) since Sun 2023-01-29 15:30:28 IST; 1 day 1h ago
     Docs: man:munged(8)
     Main PID: 1196 (munged)
        Tasks: 4
       Memory: 160.0K
      CGroup: /system.slice/munge.service
              └─1196 /usr/sbin/munged

Jan 29 15:30:18 master1 systemd[1]: Starting MUNGE authentication service...
Jan 29 15:30:28 master1 systemd[1]: Started MUNGE authentication service.
[root@master1 ~]#
```

➤ On Client:-

```
#yum --installroot=$CHROOT install -y mariadb-server mariadb-devel epel-release munge munge-libs munge-devel
rpm-build python3 perl-ExtUtils-Install gcc
```

```
#scp /etc/munge/munge.key /install/netboot/hpca2os7.9/x86_64/compute/rootimg/etc/munge
```

```
#yum --installroot=$CHROOT install -y openssl openssl-devel pam-devel numactl numactl-devel hwloc hwloc-
devel lua lua-devel readline-devel rrdtool-devel ncurses-devel man2html libibmad libibumad -y
```

```
#cd /etc/munge/
```

```
#chown munge:munge munge.key
```

```
#cd /root/rpmbuild/RPMS/x86_64/
```

```
# yum --installroot=$CHROOT install slurm*
```

```
#scp /etc/slurm/cgroup.conf $CHROOT/etc/slurm
```

```
#sinfo
```

```
[root@master1 ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug*      up    infinite     1  idle* node1
debug*      up    infinite     1  unk*  node2
[root@master1 ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug*      up    infinite     1  idle* node1
debug*      up    infinite     1  unk*  node2
[root@master1 ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug*      up    infinite     2  idle node[1-2]
[root@master1 ~]# sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
debug*      up    infinite     2  idle node[1-2]
[root@master1 ~]# █
```

```
#srun -N2 -pty /bin/bash
```

```
[root@master1 ~]# srun -N2 --pty /bin/bash
srun: job 8 queued and waiting for resources
```

```
#squeue
```

```
root@node2 ~]# packet_write_wait: connection to 192.168.168.20 port 22: broken pipe
root@master1 ~]# squeue
  JOBID PARTITION  NAME   USER ST      TIME  NODES NODELIST(REASON)
     4      debug   bash   root  R       0:27     2 node[1-2]
root@master1 ~]#
```

Ganglia Configuration

➤ On Master Node:

```
#yum install epel-release
```

```
# yum install ganglia rrdtool ganglia-gmetad ganglia-gmond ganglia-web
```

```
#htpasswd -c /etc/httpd/auth.basic admin@ganglia
```

```
# vi /etc/httpd/conf.d/ganglia.conf
```

```
#vi /etc/ganglia/gmetad.conf
```

```
data_source "my cluster" localhost

#
# Round-Robin Archives
# You can specify custom Round-Robin archives here (defaults are listed below)
#
# Old Default RRA: Keep 1 hour of metrics at 15 second resolution. 1 day at 6 minute
# RRAs "RRA:AVERAGE:0.5:1:244" "RRA:AVERAGE:0.5:24:244" "RRA:AVERAGE:0.5:168:244" "RRA:AVERAGE:0.5:672:244" \
#      "RRA:AVERAGE:0.5:5760:374"
# New Default RRA
# Keep 5856 data points at 15 second resolution assuming 15 second (default) polling. That's 1 day
# Two weeks of data points at 1 minute resolution (average)
#RRAs "RRA:AVERAGE:0.5:1:5856" "RRA:AVERAGE:0.5:4:20160" "RRA:AVERAGE:0.5:40:52704"

#
#-----
# Scalability mode. If on, we summarize over downstream grids, and respect
# authority tags. If off, we take on 2.5.0-era behavior: we do not wrap our output
# in <GRID> tags, we ignore all <GRID> tags we see, and always assume
# we are the "authority" on data source feeds. This approach does not scale to
# large groups of clusters, but is provided for backwards compatibility.
# default: on
# scalable off
#
#-----
# The name of this Grid. All the data sources above will be wrapped in a GRID
# tag with this name.
# default: unspecified
#gridname "Home office"
data_source "oxygen" 60 192.168.138.128:8649 # Master node
data_source "oxygen" 60 192.168.100.17 # Monitored node
data_source "oxygen" 60 192.168.100.20 # Monitored node
#-----
# The authority URL for this grid. Used by other gmetads to locate graphs
# for our data sources. Generally points to a ganglia/
# website on this machine.
# default: "http://hostname/ganglia/".
# where hostname is the name of this machine, as defined by gethostname().
```

#vi /etc/ganglia/gmond.conf

```
deaf = no
allow_extra_data = yes
host_dmax = 86400 /*secs. Expires (removes from web interface) hosts in 1 day */
host_tmax = 20 /*secs */
cleanup_threshold = 300 /*secs */
gexec = no
# By default gmond will use reverse DNS resolution when displaying your hostname
# Uncommenting following value will override that value.
# override_hostname = "mywebserver.domain.com"
# If you are not using multicast this value should be set to something other than 0.
# Otherwise if you restart aggregator gmond you will get empty graphs. 60 seconds is reasonable
send_metadata_interval = 0 /*secs */
}

/*
 * The cluster attributes specified will be used as part of the <CLUSTER>
 * tag that will wrap all hosts collected by this instance.
 */
cluster {
    name = "oxygen"
    owner = "unspecified"
    latlong = "unspecified"
    url = "unspecified"
}

/* The host section describes attributes of the host, like the location */
host {
    location = "unspecified"
}

/* Feel free to specify as many udp_send_channels as you like. Gmond
   used to only support having a single channel */
udp_send_channel {
    #bind_hostname = yes # Highly recommended, soon to be default.
    # This option tells gmond to use a source address
    # that resolves to the machine's hostname. Without
    # this, the metrics may appear to come from any
    # interface and the DNS names associated with
    # those IPs will be used to create the RRDs.
```

#setsebool -P httpd_can_network_connect 3

#systemctl restart httpd gmetad gmond

#systemctl enable httpd gmetad httpd

#systemctl status httpd gmetad httpd

```
[root@master1 ~]# systemctl status gmetad.service
● gmetad.service - Ganglia Meta Daemon
   Loaded: loaded (/usr/lib/systemd/system/gmetad.service; enabled; vendor preset: disabled)
   Active: active (running) since Sun 2023-01-29 18:45:31 IST; 16h ago
   Main PID: 19045 (gmetad)
     Tasks: 10
    CGroup: /system.slice/gmetad.service
            └─19045 /usr/sbin/gmetad -d 1

Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /mem buffers.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /mem buffers.rrd:...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /load fifteen.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /load fifteen.rrd:...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /cpu nice.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /cpu nice.rrd: ill...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /pkts in.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /pkts in.rrd: illeg...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /cpu idle.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /cpu idle.rrd: ill...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /disk total.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /disk total.rrd:...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /cpu system.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /cpu system.rrd:...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /mem cached.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /mem cached.rrd:...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /mem shared.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /mem shared.rrd:...second step)
Jan 30 11:11:33 master1 gmetad[19045]: RRD update (/var/lib/ganglia/rrds/oxygen/ SummaryInfo /proc_run.rrd): /var/lib/ganglia/rrds/oxygen/ SummaryInfo /proc_run.rrd: ill...second step)
Hint: Some lines were ellipsized, use -l to show in full.
[root@master1 ~]#
```

➤ **On Client Node:**

```
#export CHROOT=/install/netboot/hpcs2os7.9/x86_64/compute/rootimg
#yum --installroot=$CHROOT install ganglia rrdtool ganglia-gmetad ganglia-gmond ganglia-web
#vi /etc/ganglia/gmond.conf
```

```
deaf = no
allow_extra_data = yes
host_dmax = 86400 /*secs. Expires (removes from web interface) hosts in 1 day */
host_tmax = 20 /*secs */
cleanup_threshold = 300 /*secs */
gexec = no
# By default gmond will use reverse DNS resolution when displaying your hostname
# Uncommenting following value will override that value.
# override_hostname = "mywebserver.domain.com"
# If you are not using multicast this value should be set to something other than 0.
# Otherwise if you restart aggregator gmond you will get empty graphs. 60 seconds is reasonable
send_metadata_interval = 0 /*secs */
}

/*
 * The cluster attributes specified will be used as part of the <CLUSTER>
 * tag that will wrap all hosts collected by this instance.
 */
cluster {
    name = "oxygen"
    owner = "unspecified"
    latlong = "unspecified"
    url = "unspecified"
}

/* The host section describes attributes of the host, like the location */
host {
    location = "unspecified"
}

/* Feel free to specify as many udp_send_channels as you like. Gmond
   used to only support having a single channel */
udp_send_channel {
    #bind_hostname = yes # Highly recommended, soon to be default.
                        # This option tells gmond to use a source address
                        # that resolves to the machine's hostname. Without
                        # this, the metrics may appear to come from any
                        # interface and the DNS names associated with
                        # those IPs will be used to create the RRDs.
}
```

```
# packimage hpcs2os7.9-x86_64-netboot-compute
```

```
#systemctl status gmond
```

➤ **On Client node1 & node2:-**


```
CentOS Linux 7 (Core)
Kernel 3.10.0-1160.el7.x86_64 on an x86_64


node1 login: root
Password:
[root@node1 ~]# systemctl status gmond
Unit gmond.service could not be found.
[root@node1 ~]# systemctl status gmond
■ gmond.service - Ganglia Monitoring Daemon
   Loaded: loaded (/usr/lib/systemd/system/gmond.service; enabled; vendor preset: disabled)
   Active: active (running) since Sun 2023-01-29 18:08:37 IST; 24s ago
     Process: 1569 ExecStart=/usr/sbin/gmond (code=exited, status=0/SUCCESS)
    Main PID: 1570 (gmond)
       Tasks: 2
      Memory: 956.0K
      CGroup: /system.slice/gmond.service
              └─1570 /usr/sbin/gmond


Jan 29 18:08:37 node1.xcat.in systemd[1]: Starting Ganglia Monitoring Daemon...
Jan 29 18:08:37 node1.xcat.in systemd[1]: Started Ganglia Monitoring Daemon.
[root@node1 ~]# _
```

```
CentOS Linux 7 (Core)
Kernel 3.10.0-1160.el7.x86_64 on an x86_64

node2 login: root
Password:
[root@node2 ~]# systemctl status gmond
■ gmond.service - Ganglia Monitoring Daemon
   Loaded: loaded (/usr/lib/systemd/system/gmond.service; enabled; vendor preset: disabled)
   Active: active (running) since Mon 2023-01-30 10:42:59 IST; 38min ago
     Process: 1568 ExecStart=/usr/sbin/gmond (code=exited, status=0/SUCCESS)
    Main PID: 1569 (gmond)
       Tasks: 2
      Memory: 1.3M
      CGroup: /system.slice/gmond.service
              └─1569 /usr/sbin/gmond

Jan 30 10:42:59 node2.xcat.in systemd[1]: Starting Ganglia Monitoring Daemon...
Jan 30 10:42:59 node2.xcat.in systemd[1]: Started Ganglia Monitoring Daemon.
[root@node2 ~]#
```





Benchmarking

```
#yum install blas
# yum install epel-release
# yum install atlas
# yum install blas-devel
# wget https://netlib.org/benchmark/hpl/hpl-2.3.tar.gz
# tar -xvf hpl-2.3.tar.gz
# wget https://download.open-mpi.org/release/open-mpi/v4.1/openmpi-4.1.4.tar.gz
# tar -xvf openmpi-4.1.4.tar.gz
# cd openmpi-4.1.4
#./configure --prefix=/opt/openmpi-4.1.4 --enable-orterun-prefix-by-default
# make -j4
# make install
# export PATH=/opt/openmpi-4.1.4/bin:$PATH
# echo$PATH
# export LD_LIBRARY_PATH=/opt/openmpi-4.1.4/lib:$LD_LIBRARY_PATH
# cd /root/hpl-2.3/
# cd setup/
# cp Make.Linux_PII_CBLAS /root/hpl-2.3
#cd /root/hpl-2.3/
#vi Make.Linux_PII_CBLAS
```

```

ABCH      = Linux_FTI_CSAS
#
# -----
# - HPL Directory Structure / HPL Library -----
#
#
# TOPdir      = /root/Downloads/hpl-2.3
# INCdir      = $(TOPdir)/include
# LIBdir      = $(TOPdir)/lib/$(ARCH)
# LIBdir      = $(TOPdir)/lib/$(ARCH)
# HPLlib      = $(LIBdir)/libhpl.a
#
# -----
# - Message Passing Library (MPI) -----
#
# MPinc tells the C compiler where to find the Message Passing Library
# header files, MPlib is defined to be the name of the library to be
# used. The variable MPdir is only used for defining MPinc and MPlib.
#
# MPdir       = /opt/openmpi-4.1.4
# MPinc       = -I$(MPdir)/include
# MPlib       = $(MPdir)/lib/libmpi.so.40
#
# -----
# - Linear Algebra Library (BLAS or VSL) -----
#
# Linc tells the C compiler where to find the Linear Algebra library
# header files, LAlib is defined to be the name of the library to be
# used. The variable Ladir is only used for defining Linc and LAlib.
#
# Ladir       = /usr/lib64/atlas
# Linc        =
# LAlib       = $(Ladir)/libblas.so.3 $(Ladir)/libatlas.so.3
#
# -----
# - FTN / C Interface -----
#
# You can skip this section if and only if you are not planning to use
# a BLAS library featuring a Fortran 77 interface. Otherwise, it is
# necessary to fill out the FCDEFs variable with the appropriate
# options. "One and only one" option should be chosen in "each" of
# the 3 following categories:
#
# 1) name space (How C calls a Fortran 77 routine)
#
# -Dadd_      : all lower case and a suffixed underscore (Sun,
#               Intel, ...), [default]
# -DnoChange  : all lower case (IBM RS6000),
# -DUpCase    : all upper case (Cray),
# -Dadd_      : the FORTRAN compiler in use is f2c.
#
# 2) C and Fortran 77 integer mapping
#
# -DF77_INTEGER=int : Fortran 77 INTEGER is a C int, [default]
# -DF77_INTEGER=long : Fortran 77 INTEGER is a C long,
# -DF77_INTEGER=short : Fortran 77 INTEGER is a C short.
#
# 3) Fortran 77 string handling
#
# -DStringStyle : The string address is passed at the string loca-
#                  tion on the stack, and the string length is then

```

```
#vim HPL.dat
```

```

File Edit View Search Terminal Help
root@master1:~/Downloads/hpl-2.3/bin/Linux.PIL.CBLAS

HPLInpack benchmark input file
Innovative Computing Laboratory, University of Tennessee
HPL-out      output file name (if any)
0            device out (0=stdout,7=stderr,file)
1            # of problems sizes (N)
5580        Ns
1            # of NBS
128          NBS
0            PMAP process mapping (0=Row-,1=Column-major)
3            # of process grids (P x Q)
4            Ps
3            Qs
0            threshold
3            # of panel fact
0 1 2        PFACTs (0=Left, 1=CrouT, 2=Right)
0            # of recursive stopping criterium
2 4          NBRINs (>= 1)
1            # of panels in recursion
2            NDRIVs
3            # of recursive panel fact.
0 1 2        RFACTs (0=Left, 1=CrouT, 2=Right)
1            # of broadcast
0            BCASTs (0=1rg,1=1rf,2=2rg,3=2rf,4=Lng,5=LmM)
1            # of lookahead depth
0            DEPTHs (=0)
2            SWAP (0=bin-exch,1=long,2=mix)
64           swapping threshold
1            Li in (0=transposed,1=no-transposed) form
0            U in (0=transposed,1=no-transposed) form
1            Equilibration (0=no,1=yes)
8            memory alignment in double (> 0)

-- INSERT --
13,2

```

```
#mpirun --allow-run-as-root -np 12 --host master,node1,node2 ./xhpl HPL.dat
```

```
[root@master: Linux_PII_CHASS] mpirun -np 4 /xhpi HPL.dat --allow-run-as-root
mpirun has detected an attempt to run as root.
```

Running as root is "strongly" discouraged as any mistake (e.g., in defining MPIDIR) or bug can cause catastrophic damage to the OS file system, leaving your system in an unusable state.

We strongly suggest that you run mpirun as a non-root user.

You can override this protection by adding the "--allow-run-as-root option to the cmd line or by setting two environment variables in the following way: the variable OMPI_ALLOW_RUN_AS_ROOT=1 to indicate the desire to override this protection, and OMPI_ALLOW_RUN_AS_ROOT_CONFIRM=1 to confirm the choice and add one more layer of certainty that you want to do so.

We reiterate our advice against doing so – please proceed at your own risk.

```
[root@master: Linux_PII_CHASS] mpirun --allow-run-as-root -np 4 /xhpi.HPL.dat
HPL ERROR from process # 0, on line 395 of function MPI_pinfo:
*** Value of % less than 1 ***
```

HPL ERROR from process # 0, on line 621 of function MPI_pinfo:
*** illegal input in file MPI.dat. Exiting *** <<<

```
[root@master: Linux_PII_CHASS] mpirun --allow-run-as-root -np 4 /xhpi.HPL.dat
Primary job terminated normally, but 1 process returned
a non-zero exit code. Per user-direction, the job has been aborted.
-----
mpirun detected that one or more processes exited with non-zero status, thus causing
the job to be terminated. The first process to do was:
```

Process name: [/57979,1,0]
Exit code: 1

```
[root@master: Linux_PII_CHASS] mpirun --allow-run-as-root -np 4 /xhpi.HPL.dat
HPL ERROR from process # 0, on line 395 of function MPI_pinfo:
*** Value of % less than 1 ***
```

HPL ERROR from process # 0, on line 621 of function MPI_pinfo:
*** illegal input in file MPI.dat. Exiting *** <<<

```
[root@master: Linux_PII_CHASS] mpirun --allow-run-as-root -np 4 /xhpi.HPL.dat
Primary job terminated normally, but 1 process returned
a non-zero exit code. Per user-direction, the job has been aborted.
-----
mpirun detected that one or more processes exited with non-zero status, thus causing
the job to be terminated. The first process to do was:
```

Process name: [/59330,1,0]
Exit code: 1

```
[root@master: Linux_PII_CHASS] mpirun --allow-run-as-root -np 4 /xhpi.HPL.dat
HPL ERROR from process # 0, on line 395 of function MPI_pinfo:
*** Value of % less than 1 ***
```

HPL ERROR from process # 0, on line 621 of function MPI_pinfo:
*** illegal input in file MPI.dat. Exiting *** <<<

MPI-OpenBLAS project linkback between December 2, 2018
Written by A. Petit and R. Clint Whaley, Innovative Computing Laboratory, UTK
Modified by Dávid Lászlezek, Innovative Computing Laboratory, UTK
Reviewed by Dávid Lászlezek, Innovative Computing Laboratory, UTK

```
root@master1:~/Downloads/hpl-2.3/bin/Linux_PII_CBLAS
=====
||Ax-b||_oo/(eps*(||A||_oo*||X||_oo+||b||_oo)*N)= 4.04560734e-03 ..... PASSED
=====
T/V      N      NB      P      Q      Time      Gflops
-----
WR00R2L4 2500    128      2      2      0.27      3.9251e+01
HPL_pdgev() start time Mon Jan 30 14:46:32 2023
HPL_pdgev() end time Mon Jan 30 14:46:32 2023

||Ax-b||_oo/(eps*(||A||_oo*||X||_oo+||b||_oo)*N)= 3.07201492e-03 ..... PASSED
=====
T/V      N      NB      P      Q      Time      Gflops
-----
WR00R2C2 2500    128      2      2      0.27      3.8564e+01
HPL_pdgev() start time Mon Jan 30 14:46:32 2023
HPL_pdgev() end time Mon Jan 30 14:46:33 2023

||Ax-b||_oo/(eps*(||A||_oo*||X||_oo+||b||_oo)*N)= 4.04560734e-03 ..... PASSED
=====
T/V      N      NB      P      Q      Time      Gflops
-----
WR00R2C4 2500    128      2      2      0.32      3.2931e+01
HPL_pdgev() start time Mon Jan 30 14:46:33 2023
HPL_pdgev() end time Mon Jan 30 14:46:33 2023

||Ax-b||_oo/(eps*(||A||_oo*||X||_oo+||b||_oo)*N)= 3.49949761e-03 ..... PASSED
=====
T/V      N      NB      P      Q      Time      Gflops
-----
WR00R2R2 2500    128      2      2      0.25      4.1203e+01
HPL_pdgev() start time Mon Jan 30 14:46:33 2023
HPL_pdgev() end time Mon Jan 30 14:46:33 2023

||Ax-b||_oo/(eps*(||A||_oo*||X||_oo+||b||_oo)*N)= 4.04560734e-03 ..... PASSED
=====
T/V      N      NB      P      Q      Time      Gflops
-----
WR00R2R4 2500    128      2      2      0.33      3.1387e+01
HPL_pdgev() start time Mon Jan 30 14:46:33 2023
HPL_pdgev() end time Mon Jan 30 14:46:34 2023

||Ax-b||_oo/(eps*(||A||_oo*||X||_oo+||b||_oo)*N)= 4.04560734e-03 ..... PASSED
=====

Finished      18 tests with the following results:
              18 tests completed and passed residual checks,
               0 tests completed and failed residual checks,
               0 tests skipped because of illegal input values.

=====
End of Tests.
=====
[root@master1 Linux_PII_CBLAS]#
```

